EVIDENCE IT WORKS · LIVE RUN · 2026-10-22 15:08

A model alone guesses. With live data, it gets it right an cites.

12 recent YC startups (Winter 2025 → Fall 2026). Same input (name + website), same 4 questions, same model (claude-sonnet-5-5). Baseline = the model answering from its own knowledge. PreDiligence = live web search (Keenable) → fact extraction → verification against memory (Aiven PostgreSQL) → answer from stored evidence.

Correct answers	Confidently wrong	Answers citing a source	Per startup (live)
15% → 92%	4 → 2	0% → 100%	69 s
7/48 vs 44/48	hallucinated answers	48/48 lines with [n]	29 sources read on aver

	MODEL ALONE				PREDILIGENCE				
STARTUP	WHAT IT DOES	HQ CITY	YC BATCH	INDUSTRY	WHAT IT DOES	HQ CITY	YC BATCH	INDUSTRY	SOURC
Acyclic Labs Agent swarms that make long-running tasks exponentially fast	?	?	?	?	✓	✓	✓	✓	25
Dydema A consumer brain-computer interface you can map to anything.	?	?	?	?	✗	✓	✓	✗	11
Nxtcure Labs The full-stack AI CRO for drug discovery.	?	?	?	?	✓	?	✓	✓	12
Polemos Industries AI-guided ballistic missiles, cheap enough to field in numbe	?	?	?	?	✓	✓	✓	✓	32
Shepherd Robotics Robots for high skilled labor powering AI infrastructure	?	?	?	?	✓	✓	✓	✓	26
The Subvocal Company Talk to your computer without talking	?	?	?	?	✓	✓	✓	✓	25
Menza The AI data analyst for consumer brands	?	?	?	?	✓	✓	✓	✓	48
DiligenceSquared AI-powered market due diligence for investment decisions	✓	✗	?	✓	✓	✓	✓	✓	10
Hyperspell Your Company Brain	✓	✓	✗	✓	✓	?	✓	✓	41
Lemma Production Monitoring for AI agents	?	?	?	?	✓	✓	✓	✓	41
Exla Building agents that reads logs and fixes issues on your beh	✗	✓	✗	✓	✓	✓	✓	✓	36
Velum Labs The OS for data quality across any stack	?	?	?	?	✓	✓	✓	✓	41

✓ correct · ✗ confidently wrong · ? abstained ("unknown"). Reference answers: each startup's YC directory record. Grading: LLM judge comparing each answer to the reference (pa accepted). All results come from live runs against the deployed product (prediligence.vercel.app → Aiven Runtime), not recorded scenarios. 2 runs where the web search returned

were excluded, and 2 startups were not scored (one run did not finish, one judge call failed). Raw answers and grades: bench/results_all.json · code: github.com/adamaass/prediligence/tree/main/bench